Mystery shop: 91 real booking attempts across Boston home services

What actually happens when an agent tries to book local home services on the web - the field evidence behind the homebase thesis.

Study run overnight Sep 14-15, 2026 in Boston/Cambridge. Rebuilt and re-delivered Sep 18. Compiled here September 23, 2026. Companion document: "Homebase: the thesis, the pressure test, and where the research landed."

1. What we did, and the ground rules

Scope: 91 final-scope customer paths - 15 each across HVAC repair, movers, cleaner switch/onboarding, appliance repair and locksmith, plus 16 handyman. Categories were your pick, chosen for high discovery share (you explicitly dropped vet, beauty and makeup as wedges because routine-relationship demand bypasses platforms).

Method: real customer journeys on real provider and marketplace sites - intake forms worked to final review, pricing and availability discovery, account creation where a real customer would need one. Every path logged: source URL, outcome taxonomy (confirmed slot vs request-form dead end vs account wall vs phone-only close), price transparency, discovery-vs-execution value moment, and time-to-answer.

Hard stops (your authorization, Sep 14 10:30 PM): throwaway emails and phone numbers allowed; real form submissions and accounts allowed; **no confirmed bookings that put a real technician on a job, and nothing that spends money.** Your depth mandate: "I don't want to know that you could have put in extra effort to actually make the customer booking real."

2. Headline numbers

91	2	16	0
real customer paths completed	ended at a strictly confirmable digital booking rail	more looked bookable but died behind ZIP / identity / phone / widget gates	bookings, dispatches, leads submitted, or dollars moved

The one-line finding: "Book now" on a local-services website usually means "callback request." A strictly confirmable, price-transparent digital rail is the rare exception, not the norm.

3. The category matrix

Category	Paths	Confirmable slot	Price transparency	Dominant failure modes	Value moment
HVAC repair	15	1 yes / 9 no / 5 unknown	9 none, 1 partial, 1 full	site stalls, phone-only close, request-form dead ends	14 execution / 1 discovery
Movers	15	0 yes / 14 no / 1 unknown	5 full, 5 partial, 5 none	quote-request orchestration (8 of 15), callbacks, estimate forms	15 discovery / 0 execution
Cleaner switch/onboarding	15	1 yes / 9 no / 5 unknown	3 full, 1 partial, 7 none	location gates, mandatory-phone leads, captcha blocks	9 execution / 6 discovery
Appliance repair	15	0 yes / 11 no / 4 unknown	1 full, 2 partial, 11 none	site stalls, captcha blocks, phone/email only, account walls	14 execution / 1 discovery
Handyman	16	0 yes / 11 no / 5 unknown	3 full, 3 partial, 10 none	site stalls, mandatory-phone leads, account/scope gates	10 discovery / 6 execution
Locksmith	15	0 yes / 14 no / 1 unknown	2 full, 4 partial, 9 none	phone-only close (8 of 15), mandatory-phone leads, phone-first	15 execution / 0 discovery

Median time-to-answer was 2-3 minutes per path across categories - the journeys are fast to fail, which is itself the finding: the funnel collapses early, at the exact moment an agent would need structured data.

4. Category-by-category

4.1 Cleaner switch/onboarding - the best first wedge

The only category where an operator (Umess, on BookingKoala) exposed a full **price -> date -> summary -> payment** rail in the inspected state. Phone-gated competitors leave room. Cleaner demand is episodic at exactly the right moments: 70-90% new-provider share at move/switch/onboarding (hypothesis) even though steady-state rebooking is relationship-owned.

4.2 HVAC repair - the best second

Urgency plus high ticket. Larger operators run ServiceTitan/Square rails, so scheduling is more credible than elsewhere, but remote price stays a range and intake forms (ownership, contact, system details) gate everything. One path did reach a confirmable slot - the only one outside cleaning.

4.3 Movers - a discovery wedge, not a booking wedge

Pure quote orchestration: zero instant-book across 15 paths, with 8 ending in quote requests. Price transparency is actually the best of the six (10 of 15 full or partial - e.g. Boston Best Rate publishes $149/hr), but nothing is bookable online. The product here is comparison and orchestration, not instant execution.

4.4 Handyman - demand is real, scope is the problem

The most discovery-heavy mix (10 of 16) and the most marketplace-mediated (8 of 16 via marketplaces). The blocker is representing scope/quality before service, not scheduling.

4.5 Appliance repair - needs diagnostic standardization first

Zero confirmable rails, worst price opacity (11 of 15 none). Same-day is widely advertised but the scheduler needs the diagnosis before the slot, which is a standardization problem an agent could own.

4.6 Locksmith - a voice-agent wedge, not a web wedge

93% of paths (14 of 15) hit a phone/callback wall. For this category **a phone call is effectively the API** - the winning product is a voice agent, not a web integration.

5. Cross-cutting findings (the ones that should shape the product)

1. A phone number is infrastructure, not an edge feature. Most funnels demand a valid mobile before showing any price - sometimes before even a planning range. A buyer-side agent without a calling/SMS identity cannot transact in these categories. This is also why zero leads were submitted even after you lifted the submission restriction: every submit-ready form required a valid phone and no throwaway phone route existed, and we never used yours. Provider response-latency is marked **unmeasured** rather than estimated.

2. ServiceTitan, Square and BookingKoala appear exactly where booking is most real. The credible-scheduling paths in the study ran on those rails. Implication: integrate at those rails, and bridge the long tail by voice.

3. The funnel fails before the transaction, not at it. Failure modes cluster at intake: location gates, identity walls, mandatory phone, captcha/bot blocks, "Book Now" buttons that open contact forms. The last mile (payment, confirmation) barely got exercised because almost nothing reached it.

6. Live case study: the plumber hunt (Sep 14, 1 Western Ave)

While the research ran, you asked for a real plumber booking for the next day - home-services intent #1 of the dataset. What the hunt showed in practice:

- **128 Plumbing:** $89 tech visit, four real windows on the next day's booking page - the strongest option, and even here the slot was "requested, not guaranteed."
- **Winters Home Services:** $29 estimate but callback-only - cheap entry, no commit.
- **Roto-Rooter:** 24/7, drain/sewer focus - emergency posture, phone-led.
- **Taskrabbit:** ~$77/hr, only sensible for simple fixture jobs.

You pushed back on paying $89 over $29 and asked the agent to call - it can't make calls, so both numbers went back to you. The job details, window and go/no-go were never supplied, so nothing was booked. The episode is a tidy demonstration of the thesis: the agent found and compared supply fine, then hit the exact wall the research predicts - commitment requires a phone identity and a guarantee layer nobody offers.

7. Limits of the study

- Boston/Cambridge only; ~15 paths per category is directional, not statistical.
- "Unknown" confirmability (5-6 paths in several categories) reflects gates that couldn't be passed without a real phone identity - itself a finding, but it caps precision.
- No lead submissions means provider response speed and sales aggressiveness are unmeasured.
- Paths were run by an agent that sites can detect as automation; captcha/bot blocks (several paths) may overstate friction vs. a human - or accurately reflect the friction an agent will face in production. Both readings matter for the thesis.

8. Next experiments (proposed, not yet run)

1. **25 cleaner-switch + 25 HVAC intents** with the throwaway-phone problem solved - measures response latency, follow-up persistence, and real confirm rates.
2. **Kernel falsification test:** does 70%+ of the request -> hold -> confirm state machine transfer across both categories? If yes, the execution layer is horizontal even if the wedge is vertical.
3. **Voice-agent probe for locksmith:** the one category where the phone is the API - test whether an AI voice front can complete what the web cannot.